

REPULABS · FEATURE ROADMAP

12 candidate features for the next 4 months.

An honest, dev-ready inventory of what to build, why each one matters, what it requires, and where it fits in a phased rollout.

v1.0 · May 18, 2026

Prepared for the Repulabs founding team

Contents

#	Feature	Phase
01	NFC Tap Products (Plaques, Cards, Stickers)	Phase 2
02	Airbnb Reviews Integration	Phase 3
03	WiFi NFC Cards	Phase 4
04	AI Phone Receptionist + Calendar Booking	Phase 1 add-on
05	AI Reputation Forecaster	Phase 6
06	Phone-Call Sentiment Early-Warning	Phase 6
07	Review-to-Marketing Asset Pipeline	Phase 5
08	AI Reputation Coach — Monday Standup	Phase 5
09	Competitor Pulse + AI Mystery Shopper	Phase 7
10	Multi-Channel Inbox AI Triage	Phase 6
11	Photo Review Booster	Phase 6
12	Voice Survey via Phone	Phase 7
	Phased rollout summary	Page 14

How to read this document. Each feature has the same shape — *Why it matters*, *What it requires*, and a stat row showing build effort + business impact. Features are ordered by rollout phase, not by importance. The phased summary at the end maps features to weeks of calendar time.

01

NFC Tap Products (Plaques, Cards, Stickers)

Hardware product line · Phase 2

Why it matters.

QR scanning requires opening a camera, framing the code, and tapping a notification. NFC is one motion: tap, done. We project a 35–45% lift in scan-to-review conversion purely from removing camera friction, especially for older customers and rushed lunch crowds. Same backend, same activation flow — the chip is a delivery mechanism, not a new feature surface.

What it requires.

NTAG216 chip supplier; NDEF programmer (~\$50); minor schema migration (product_kind, nfc_chip_uid).

EFFORT	IMPACT
5 days build + 2–4 weeks factory pilot.	New \$89 SKU tier above the current \$59.99 plaque. Margin: 60–70%.

02

Airbnb Reviews Integration

Channel expansion · Phase 3

Why it matters.

Hospitality is one of the largest verticals for reputation management and Airbnb is the single biggest review source for short-term-rental operators. Airbnb has no public API, so this is a genuine moat — competitors avoid it because parsing emails feels gross. Done well, it unlocks an entire vertical (300K+ Airbnb hosts in US/AU/UK alone).

What it requires.

Resend inbound DNS configuration; email template parser; new external_review_id column.

EFFORT	IMPACT
5 days.	Opens hospitality vertical; expected 25% TAM expansion.

03

WiFi NFC Cards

Hardware product line · Phase 4

Why it matters.

Adjacent SKU with zero backend complexity — customer programs the chip themselves using a free phone app. Sells as an accessory at the counter, gateway product for businesses not yet ready for review management. Drives top-of-funnel + brand familiarity.

What it requires.

NFC card SKU artwork; help page; supplier; 1-page printed insert.

EFFORT	IMPACT
1 day.	Low-effort revenue diversifier. \$19–29 SKU.

04 AI Phone Receptionist + Calendar Booking

Already 80% built · Phase 1 add-on

Why it matters.

The full pipeline exists: Twilio webhooks, Claude Haiku dialog brain, ElevenLabs TTS, Cal.com integration for slot lookup and booking. Missing: branded confirmation emails to both the caller and the business owner (~6 hours of glue code). Closing this gap turns an internal feature into a customer-visible differentiator.

What it requires.

Resend transactional templates; PhoneBooking.notifiedAt column; per-org sender domain (already plumbed for outreach).

EFFORT	IMPACT
6 hours.	Activates an existing Pro tier feature; reduces support load on "did we get the booking?" questions.

05 AI Reputation Forecaster

Predictive analytics · Phase 6

Why it matters.

Every competitor shows past metrics. None show the future one. Predicting the next 90 days of star-rating trajectory with confidence bands and lever toggles ("reply to all pending in 48h → +0.06") turns reputation from a lagging indicator into a leading one. Massive retention moat — operators won't switch off the platform predicting their future.

What it requires.

ARIMA + Bayesian credible-interval model running nightly; new reputation_forecasts table.

EFFORT	IMPACT
6 days build, 2 weeks of customer data before predictions stabilize.	Headline marketing feature; reduces churn.

06 Phone-Call Sentiment Early-Warning

Real-time intervention · Phase 6

Why it matters.

Industry standard is "respond within 24 hours of a 1-star review." Our standard would be "intercept the upset caller before they leave one." Sentiment scoring on existing PhoneCallTurn rows + SMS-the-manager-now + auto-drafted recovery voucher = a measurable drop in 1-star reviews. Only possible because we own the phone channel; competitors don't.

What it requires.

Sentiment-classification tool call in brain.ts; new phone_call_escalations table; Twilio live-listen conference mode.

EFFORT	IMPACT
5 days.	Customer-recovery win-rate improvement; testimonial-grade feature.

07 Review-to-Marketing Asset Pipeline

Closed-loop value · Phase 5

Why it matters.

Competitors stop at "got the review." We'd convert each 5-star review into six marketing assets in 30 seconds: Instagram square, LinkedIn carousel slide, embeddable website testimonial widget, Google Ads RSA variants, 1-page case study PDF, and an email signature trust badge. Closes the loop from

review → asset → distribution → more customers.

What it requires.

next/og image rendering; Claude generation prompts; reuse of existing Meta/LinkedIn OAuth.

EFFORT	IMPACT
7 days.	Sellable as standalone Pro feature; visible ROI per review.

08 AI Reputation Coach — Monday Standup

Retention + engagement · Phase 5

Why it matters.

Replaces a fractional reputation manager (\$800–2,000/month elsewhere). Weekly Monday voice memo + SMS + email synthesizing last week's data and prescribing this week's action. Built on top of the forecaster + topic engine + customer DB. Drives weekly platform engagement, which is the single biggest predictor of retention in SaaS.

What it requires.

Weekly cron; reputation_briefs table; ElevenLabs TTS (already integrated); SMS + email delivery (already integrated).

EFFORT	IMPACT
5 days.	Major retention lever; weekly touch keeps platform top-of-mind.

09 Competitor Pulse + AI Mystery Shopper

Premium add-on · Phase 7

Why it matters.

Every operator wants to know what the competing business across the street is doing better. Today they ask their friends; we'd answer with data. Nightly public-page scraping for review velocity + AI agent monthly calls competitors as a customer and scores them. Premium tier upsell driver — \$59.99 → \$99/mo.

What it requires.

Bright Data residential proxy (~\$50/mo per portfolio); outbound calling reuses existing infrastructure; HTML-resilient parser.

EFFORT	IMPACT
10 days.	Highest revenue impact per customer; clear premium-tier driver.

10 Multi-Channel Inbox AI Triage

Workflow consolidation · Phase 6

Why it matters.

Reviews + Facebook comments + Instagram DMs + SMS replies + voicemails + emails arrive in six different inboxes. We'd unify them with AI-categorized urgency + sentiment + topic, routed to the right team member. Owners stop tab-juggling; response time drops from hours to minutes.

What it requires.

Provider webhook handlers (Meta, Instagram); unified message table; AI triage policy engine.

EFFORT	IMPACT
8 days.	Daily-use platform feature; sticky workflow.

11 Photo Review Booster

Targeted outreach · Phase 6

Why it matters.

Photo-attached Google reviews carry 2× the SEO weight of text-only ones. Customers most likely to attach photos are repeat high-spenders. We'd identify them automatically from review history + order data (when CRM-connected) and craft a tailored ask. Higher quality review yield with less outreach volume.

What it requires.

Reviewer-profile enrichment; CRM-data correlation; specialized request template.

EFFORT	IMPACT
4 days.	Higher-quality review yield; better Maps ranking signal.

12 Voice Survey via Phone

After-service NPS · Phase 7

Why it matters.

AI calls a customer 24–48 hours after a service for a 60-second voice survey. Tone analysis catches genuine sentiment that text NPS misses (one customer types "great!" while another sighs and says "yeah it was fine" — the difference is everything). Premium-tier feature; charged per completed call.

What it requires.

Outbound calling already exists; sentiment + transcription pass; survey-script engine; consent recording.

EFFORT	IMPACT
6 days.	Premium add-on; new revenue stream.

Phased rollout summary

Each phase below is sized so a single founder + AI-paired engineer can ship it without blocking on the next phase. Phase 1 is already underway as of this document's date.

Phase	Scope	Effort
1 — this week	Restore-from-trash + /signup + smart /not-activated + booking-confirmation emails	~1 day
2 — next 2 weeks	NFC product line (schema, factory workflow, admin UI)	~5 days
3 — next 4 weeks	Airbnb via Resend inbound email parsing	~5 days
4 — opportunistic	WiFi NFC SKU (product + insert + help page)	~1 day
5 — month 2	Reputation Coach + Review-to-Marketing pipeline	~12 days
6 — month 3	Phone Sentiment + Reputation Forecaster + Multi-channel Inbox + Product Booster	~25 days
7 — month 4 (premium)	Competitor Pulse + Mystery Shopper + Voice Survey	~16 days

Footnotes

Effort estimates assume one full-time founder paired with AI-accelerated development. Calendar time may be longer due to factory pilot lead times (NFC, Phase 2) and OAuth verification waits (Phase 6 if integrating with Google's Business Profile sensitive scope).

Revenue impact numbers are conservative. Several features (Reputation Coach, Competitor Pulse) double as standalone SKUs and could justify a \$40–60/month price tier above the current \$59.99 Pro plan.

The roadmap is a living document. Customer feedback in the first 60 days will reorder Phase 5 and beyond — Reputation Coach and Review-to-Marketing pipeline both have signals from early users that they'd unblock close deals, but that's worth validating against real intent before committing engineering weeks.